

Yes — **absolutely**, the **Von Neumann bottleneck** (and its modern manifestations in **HPC, AI, and memory hierarchy**) is *still an open, active research field*.

You could **definitely** make this your **Master’s thesis topic**, *if you frame it well*.

Let’s break this down carefully so you know how to approach it strategically 📌

🎯 1. Why this can be a thesis topic

The Von Neumann bottleneck isn’t “old” — it has just **changed form**.

In HPC, it shows up as:

- **Memory bandwidth limits** (CPU or GPU waiting for RAM)
- **Data movement energy costs** (more power to move data than compute)
- **Communication bottlenecks in distributed memory systems**

Modern research calls this the **Memory Wall** problem.

Your thesis could aim to:

Reduce, hide, or rethink the bottleneck using modern architectures, algorithms, or data locality techniques.

🧠 2. Possible Research Directions

Here are real, *researchable paths* you can take:

Category	Focus	Example Thesis Idea
Hardware-aware computation	Optimize data movement between CPU–GPU–Memory	“Improving data locality in heterogeneous memory systems for AI workloads.”
Processing-in-Memory (PIM)	Prototype or simulate in-memory compute	“Evaluating in-memory matrix multiplication accelerators for HPC.”
Compiler-level optimization	Reordering memory access patterns	“Compiler-driven prefetch scheduling to reduce Von Neumann bottleneck in HPC kernels.”
Cache optimization	Better use of cache levels	“Adaptive caching and prefetching for large-scale scientific computation.”
Energy efficiency	Minimize data movement energy	“Quantifying energy cost of memory transfer in AI accelerators.”

Each of these ties back to your central theme — mitigating the bottleneck.

3. Steps to Turn This Into a Thesis

Step 1: Ground yourself in background

Learn about:

- **Memory hierarchy** (L1/L2/L3 cache, DRAM, HBM)
- **CPU/GPU pipelines**
- **PIM and Near-Memory Computing**
- **Roofline model** (used to analyze compute vs memory performance)

Books/papers:

- “*Computer Architecture: A Quantitative Approach*” – Hennessy & Patterson
 - Research papers on “Memory wall” and “Processing-in-Memory”
-

Step 2: Define your scope

Instead of “solving” the bottleneck (too big), focus on:

“Reducing the bottleneck for [specific workload or architecture] using [specific technique].”

e.g.,

- “*Reducing data movement in stencil computations using PIM simulation.*”
 - “*Quantifying performance gain using compiler-based prefetching in HPC.*”
-

Step 3: Talk to researchers

Yes  — **interviewing researchers** is a very good move.

Here's how:

- Find professors or PhDs publishing in *memory systems*, *HPC architecture*, or *PIM*.
- Email them (even short questions).
- Ask about open problems, bottlenecks they face in HPC codes, or simulation tools they use.

You might even get invited to collaborate or co-author a paper if your direction is solid.

Step 4: Get hands-on

You can:

- Use **gem5**, **Sniper**, or **SimpleScalar** simulators for CPU/memory modeling.
- Write small C/OpenMP programs and measure cache miss rates with tools like **perf** or **PAPI**.
- Compare performance with different memory access patterns.

This gives your thesis *quantitative results*.

Step 5: Build your network

Attend small **HPC seminars** or **IEEE/ACM events** online.

Talk to people in:

- **Memory systems**
- **Computer architecture**
- **Parallel computing**

That's where you'll find mentors who can guide your research scope.



4. The Vision

You could name your research something like:

"Towards Breaking the Von Neumann Bottleneck in HPC: Exploring Memory-Centric Architectures and Dataflow Models"

That alone sounds like a *Master's thesis worth reading*.